



# mongoDB

## Sharding

Prof. P. M. Jadav  
Associate Professor  
Computer Engineering Department  
Faculty of Technology  
D. D. University, Nadiad

# Introduction

- **Sharding** is the process of **splitting data across multiple machines**, known as **shards**, to:
  - Handle large datasets
  - Ensure high throughput (read/write performance)
  - Scale horizontally (add more servers to distribute load)

# Understanding Partitioning

Original Table

CUSTOMER ID	FIRST NAME	LAST NAME	FAVORITE COLOR
1	TAEKO	OHNUKI	BLUE
2	O.V.	WRIGHT	GREEN
3	SELDA	BAĞCAN	PURPLE
4	JIM	PEPPER	AUBERGINE

Vertical Partitions

VP1

CUSTOMER ID	FIRST NAME	LAST NAME
1	TAEKO	OHNUKI
2	O.V.	WRIGHT
3	SELDA	BAĞCAN
4	JIM	PEPPER

VP2

CUSTOMER ID	FAVORITE COLOR
1	BLUE
2	GREEN
3	PURPLE
4	AUBERGINE

Horizontal Partitions

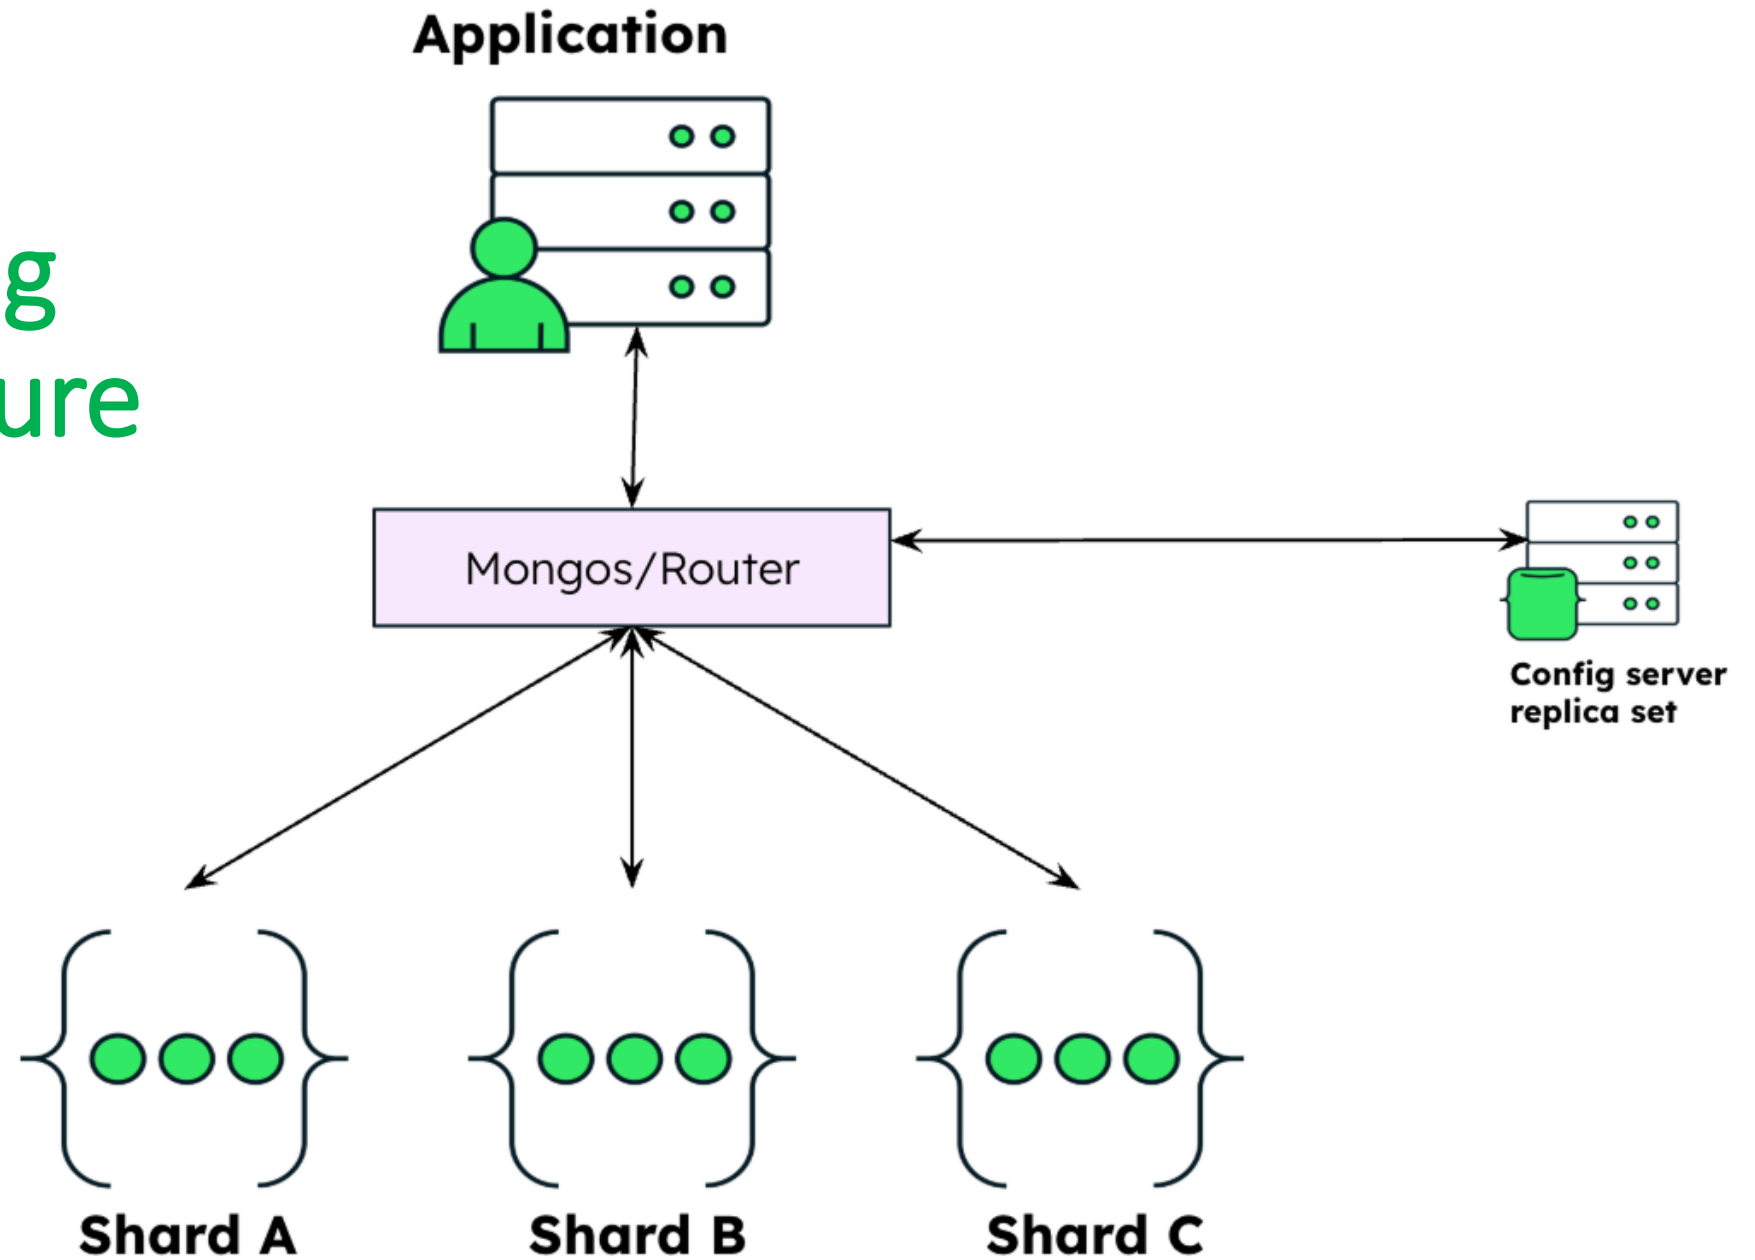
HP1

CUSTOMER ID	FIRST NAME	LAST NAME	FAVORITE COLOR
1	TAEKO	OHNUKI	BLUE
2	O.V.	WRIGHT	GREEN

HP2

CUSTOMER ID	FIRST NAME	LAST NAME	FAVORITE COLOR
3	SELDA	BAĞCAN	PURPLE
4	JIM	PEPPER	AUBERGINE

# Sharding Architecture



# Sharding Components

## 1. Shard

- Each **shard** holds a portion of the data
- It contains a subset of the cluster's data
- Each shard has a primary and two secondary nodes by default
- The data is **partitioned** among shards based on a **shard key**

# Sharding Components

## 2. Shard Key

- A field(s) used to determine the distribution of data across shards
- Must be chosen carefully as it impacts **performance, query distribution, and data balancing**
- Types:
  - **Ranged Sharding** – based on value ranges of the shard key
  - **Hashed Sharding** – uses a hash of the shard key for uniform distribution

# Sharding Components

## 3. Config Server Replica Set (CSRS)

- Store the **metadata** about the cluster and data distribution
- Minimum **3 config servers** are required in production for fault tolerance
- They maintain the mapping between shard key ranges and shards

# Sharding Components

## 4. Query Routers (mongos)

- Mongos instances route queries from clients to the appropriate shard(s)
- The client application connects to mongos, not directly to the shards
- mongos consults config servers to locate the relevant data



# Advantages

- 1) **Increased read/write throughput** — By distributing the dataset across multiple shards, read/write operation capacity is increased as long as read and write operations are confined to a single shard
- 2) **Increased storage capacity** — Similarly, by increasing the number of shards, you can also increase overall total storage capacity
- 3) **High availability** — Each shard is a replica set, every piece of data is replicated. Even if an entire shard becomes unavailable since the data is distributed, the database as a whole still remains partially functional, with part of the schema on different shards

# Disadvantages

## 1) Query overhead

- Each sharded database must have a **separate machine** or **service** which understands how **to route a querying operation** to the appropriate shard
- This introduces additional **latency** on every operation
- If the data required for the query is horizontally partitioned across multiple shards, the **router** must then query each shard and **merge the result** together
- This can make an otherwise **simple operation** quite **expensive** and slow down response times

# Disadvantages

## 2) Complexity of administration

- With every sharded database, on top of managing the shards themselves, there are **additional service nodes** to maintain
- In cases where replication is being used, any **data updates must be mirrored** across each replicated node
- Overall, a sharded database is a **more complex** system which requires **more administration**

# Disadvantages

## 3) Increased infrastructure costs

- Sharding by its nature requires **additional machines** and compute power over a single database server
- While this allows your database to grow beyond the limits of a single machine, each additional shard comes with **higher costs**
- The cost of a distributed database system, especially if it is missing the **proper optimization**, can be significant

# Sharding Types

- 1) Range/Dynamic Sharding
- 2) Algorithmic/Hashed Sharding
- 3) Entity-Relationship based Sharding
- 4) Geography based Sharding

# Range/ Dynamic Sharding

PRODUCT	PRICE
WIDGET	\$118
GIZMO	\$88
TRINKET	\$37
THINGAMAJIG	\$18
DOODAD	\$60
TCHOTCHKE	\$999



(\$0-\$49.99)

PRODUCT	PRICE
TRINKET	\$37
THINGAMAJIG	\$18



(\$50-\$99.99)

PRODUCT	PRICE
GIZMO	\$88
DOODAD	\$60



(\$100+)

PRODUCT	PRICE
WIDGET	\$118
TCHOTCHKE	\$999

# Range/Dynamic Sharding

- It takes a field on the record as an input and, based on a predefined range, allocates that record to the appropriate shard
- The field on which the range is based is also known as the [shard key](#)
- A poor choice of shard key will lead to unbalanced shards, which leads to decreased performance
- An effective shard key will allow for queries to be targeted to a minimum number of shards

# Algorithmic/Hashed Sharding

- It takes a record as an input and applies a hash function or algorithm to it which generates an output or hash value
- This output is then used to allocate each record to the appropriate shard
- Hashed sharding typically disregards the meaning of the data



# Algorithmic/Hashed Sharding

**Shard Key**

COLUMN 1	COLUMN 2	COLUMN 3
A		
B		
C		
D		



**HASH  
FUNCTION**



COLUMN 1	HASH VALUES
A	1
B	2
C	1
D	2



**Shard 1**

COLUMN 1	COLUMN 2	COLUMN 3
A		
C		

**Shard 2**

COLUMN 1	COLUMN 2	COLUMN 3
B		
D		

# Algorithmic/Hashed Sharding

- Advantages:

- Hashing the inputs allows more even distribution across shards even when there is not a suitable shard key
- No lookup table needs to be maintained

- Disadvantages:

- Query operations for multiple records are more likely to get distributed across multiple shards
- Resharding can be expensive
- Any update to the number of shards likely requires rebalancing all shards to moving around records

# Entity-Relationship based Sharding

- It keeps related data together on a single physical shard
- In a relational database related data is often spread across several different tables
- e.g. Consider the case of a shopping database with users and payment methods
  - Each user has a set of payment methods that is tied tightly with that user
  - Keeping related data together on the same shard can reduce the need for broadcast operations, increasing performance

# Geography based Sharding

- It keeps related data together on a single shard
- The data is related by geography
- This is essentially ranged sharding where the shard key contains geographic information and the shards themselves are geo-located
- For example, consider a dataset where each record contains a “country” field
- In this case, we can both increase overall performance and decrease system latency by creating a shard for each country or region, and storing the appropriate data on that shard

# References

- <https://www.mongodb.com/docs/manual/sharding/>
- <https://www.mongodb.com/resources/products/capabilities/database-sharding-explained>
- <https://www.mongodb.com/resources/products/capabilities/sharding>
- <https://www.digitalocean.com/community/tutorials/understanding-database-sharding>